

NAME: SURANA GAURAV

ROLL: 20BEE0371

ASSIGNMENT 3

Date : 06-06-2023

```
In [3]: import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
import os
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import LabelEncoder
```

```
df = pd.read_csv('Housing.csv')
df.head()
```

```
Out[3]:
```


	price	area	bedrooms	bathrooms	stories	mainroad	guestroom	basement	hotwaterheating	airconditioning	parking	furnishingstatus
0	13300000	7420	4	2	3	yes	no	no	no	yes	2	furnished
1	12250000	8060	4	4	4	yes	no	no	no	yes	3	furnished
2	12250000	9960	3	2	2	yes	no	yes	no	no	2	semi-furnished
3	12215000	7500	4	2	2	yes	no	yes	no	yes	3	furnished
4	11410000	7420	4	1	2	yes	yes	yes	no	yes	2	furnished

```
In [88]: df.dtypes
```


```
Out[88]: price          float64
area          float64
bedrooms      int64
bathrooms     int64
stories       int64
mainroad      int64
guestroom     int64
basement      object
hotwaterheating int64
airconditioning int64
parking       float64
furnishingstatus object
dtype: object
```

```
In [4]: # Plotting a histogram(univariate analysis)
plt.hist(df['price'], bins=10)
plt.xlabel('price')
```

```
In [7]: # Plotting a Histogram(univariate analysis)
plt.hist(df['price'], bins=10)
plt.xlabel('price')
plt.ylabel('stories')
plt.title('Histogram of price')
plt.show()
```



```
In [9]: # Plotting a box plot(univariate analysis)
plt.boxplot(df['bedrooms'])
plt.ylabel('bedrooms')
plt.title('Box Plot of Variable')
plt.show()
```



```

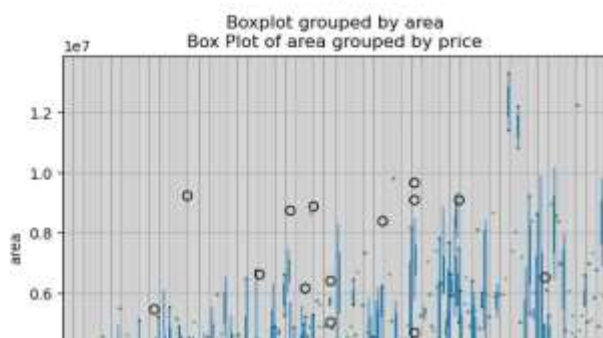


In [43]: # Scatter plot(bi-variate analysis )
plt.scatter(df['price'], df['area'])
plt.xlabel('price')
plt.ylabel('area')
plt.title('Scatter Plot of price vs area')
plt.show()

# Line plot(bi-variate analysis )
plt.plot(df['price'], df['area'])
plt.xlabel('price')
plt.ylabel('area')
plt.title('Line Plot of price vs area')
plt.show()

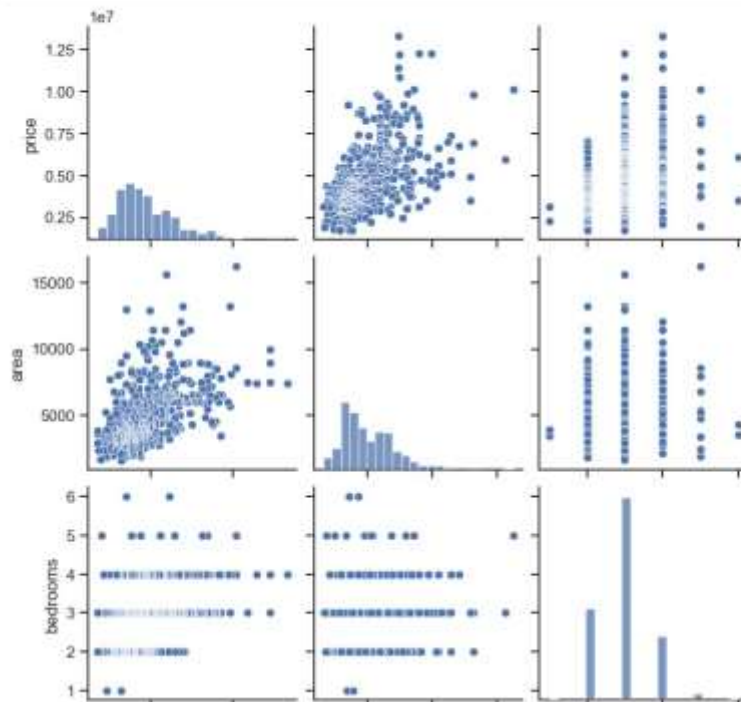
# Box plot(bi-variate analysis )
df.boxplot(column='price', by='area')
plt.xlabel('price')
plt.ylabel('area')
plt.title('Box Plot of area grouped by price')
plt.show()

# Correlation matrix(bi-variate analysis )
correlation_matrix = df[['price', 'area']].corr()
print(correlation_matrix)

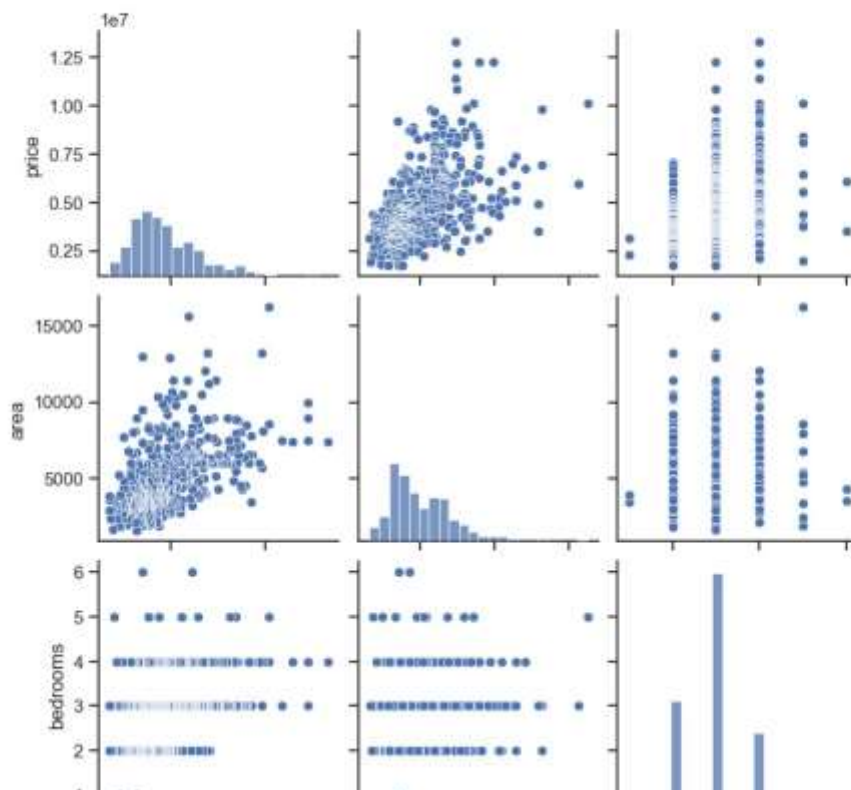
```




```
In [10]: #Pairplot(multivariate analysis)
sns.pairplot(df, vars=['price', 'area', 'bedrooms'])
plt.show()
```



```
In [11]: # Scatter plot matrix(multivariate analysis)
sns.set(style='ticks')
sns.pairplot(df, vars=['price', 'area', 'bedrooms'], kind='scatter')
plt.show()
```



```
In [17]: # Heatmap(multivariate analysis)
correlation_matrix = df[['price', 'area', 'bedrooms']].corr()
sns.heatmap(correlation_matrix, annot=True)
plt.title('Correlation Heatmap')
plt.show()
```



```
In [18]: # Get summary statistics of the entire dataset(descriptive statistics)
print(df.describe())
```

	price	area	bedrooms	bathrooms	stories	%
count	5.450000e+02	545.000000	545.000000	545.000000	545.000000	545.000000
mean	4.796729e+06	5150.541284	1.865118	1.286139	1.805595	1.805595
std	1.870440e+06	2170.141023	0.738064	0.502470	0.867401	0.867401
min	1.750000e+06	1650.000000	1.000000	1.000000	1.000000	1.000000
25%	3.430000e+06	3600.000000	2.000000	1.000000	1.000000	1.000000
50%	4.340000e+06	4500.000000	3.000000	1.000000	2.000000	2.000000
75%	5.740000e+06	6160.000000	3.000000	2.000000	2.000000	2.000000
max	1.330000e+07	16200.000000	6.000000	4.000000	4.000000	4.000000

```
In [22]: # Get summary statistics for a specific column/variable(descriptive statistics)
print(df['area'].describe())
```

count	545.000000
mean	5150.541284
std	2170.141023
min	1650.000000
25%	3600.000000
50%	4500.000000
75%	6160.000000
max	16200.000000
Name: area, dtype: float64	

```
In [26]: # Check for missing values
missing_values = df.isnull().sum()
print(missing_values)
#No missing values
```

price	0
area	0
bedrooms	0
bathrooms	0
stories	0
mainroad	0
guestroom	0
basement	0
hotwaterheating	0
airconditioning	0
parking	0
furnishingstatus	0
dtype: int64	

```
In [25]: # Defining the columns for which we want to detect outliers
columns_to_check = ['price', 'area']

# Calculate the IQR for each column
Q1 = df[columns_to_check].quantile(0.25)
Q3 = df[columns_to_check].quantile(0.75)
IQR = Q3 - Q1

# Detecting outliers based on the IQR method
outliers = ((df[columns_to_check] < (Q1 - 1.5 * IQR)) | (df[columns_to_check] > (Q3 + 1.5 * IQR)))

# Replacing outliers with the median value
df[outliers] = np.nan # Replacing outliers with NaN
df_filled = df.fillna(df.median()) # Replacing NaN with the median
```

```

In [45]: # Perform label encoding for each categorical column
label_encoder = LabelEncoder()
for column in categorical_columns:
    df[column] = label_encoder.fit_transform(df[column])

# Print the encoded dataset
print(df)

```

	price	area	bedrooms	bathrooms	stories	mainroad	guestroom	\
0	13300000	7420	4	2	3	1	0	
1	12250000	8960	4	4	4	1	0	
2	12250000	9960	3	2	2	1	0	
3	12215000	7500	4	2	2	1	0	
4	11410000	7420	4	1	2	1	1	
...	
540	1820000	3000	2	1	1	1	0	
541	1767150	2400	3	1	1	0	0	
542	1750000	3620	2	1	1	1	0	
543	1750000	2910	3	1	1	0	0	
544	1750000	3850	3	1	2	1	0	

	basement	hotwaterheating	airconditioning	parking	furnishingstatus
0	no	0	yes	2	furnished
1	no	0	yes	3	furnished
2	yes	0	no	2	semi-furnished
3	yes	0	yes	3	furnished
4	yes	0	yes	2	furnished
...
540	yes	0	no	2	unfurnished
541	no	0	no	0	semi-furnished
542	no	0	no	0	unfurnished
543	no	0	no	0	furnished
544	no	0	no	0	unfurnished

[545 rows x 12 columns]

```

In [46]: # Perform label encoding for each categorical column
label_encoder = LabelEncoder()
for column in categorical_columns:
    df[column] = label_encoder.fit_transform(df[column])

# Print the encoded dataset
print(df)

```

	price	area	bedrooms	bathrooms	stories	mainroad	guestroom	\
0	13300000	7420	4	2	3	1	0	
1	12250000	8960	4	4	4	1	0	
2	12250000	9960	3	2	2	1	0	

```

In [47]: # Perform label encoding for each categorical column
label_encoder = LabelEncoder()
for column in categorical_columns:
    df[column] = label_encoder.fit_transform(df[column])

# Print the encoded dataset
print(df)

```

	price	area	bedrooms	bathrooms	stories	mainroad	guestroom	\
0	13300000	7420	4	2	3	1	0	
1	12250000	8960	4	4	4	1	0	
2	12250000	9960	3	2	2	1	0	
3	12215000	7500	4	2	2	1	0	
4	11410000	7420	4	1	2	1	1	
...	
540	1820000	3000	2	1	1	1	0	
541	1767150	2400	3	1	1	0	0	
542	1750000	3620	2	1	1	1	0	
543	1750000	2910	3	1	1	0	0	
544	1750000	3850	3	1	2	1	0	

	basement	hotwaterheating	airconditioning	parking	furnishingstatus
0	no	0	yes	2	furnished
1	no	0	yes	3	furnished
2	yes	0	no	2	semi-furnished
3	yes	0	yes	3	furnished
4	yes	0	yes	2	furnished
...
540	yes	0	no	2	unfurnished
541	no	0	no	0	semi-furnished
542	no	0	no	0	unfurnished
543	no	0	no	0	furnished
544	no	0	no	0	unfurnished

[545 rows x 12 columns]

```

In [48]: # Perform label encoding for each categorical column
label_encoder = LabelEncoder()
for column in categorical_columns:
    df[column] = label_encoder.fit_transform(df[column])

# Print the encoded dataset
print(df)

```

	price	area	bedrooms	bathrooms	stories	mainroad	guestroom	\
0	13300000	7420	4	2	3	1	0	
1	12250000	8960	4	4	4	1	0	

```
In [5]: # Split the data into dependent and independent variables
X = df.iloc[:, :-1] # Select all columns except the last one as independent variables
y = df.iloc[:, -1] # Select the last column as the dependent variable

# Print the independent variables (X)
print(X)

# Print the dependent variable (y)
print(y)
```

	price	area	bedrooms	bathrooms	stories	mainroad	guestroom	\
0	13300000	7420	4	2	3	1	0	
1	12250000	8960	4	4	4	1	0	
2	12250000	9960	3	2	2	1	0	
3	12215000	7500	4	2	2	1	0	
4	11410000	7420	4	1	2	1	1	
...
540	1820000	3000	2	1	1	1	0	
541	1767150	2400	3	1	1	0	0	
542	1750000	3620	2	1	1	1	0	
543	1750000	2910	3	1	1	0	0	
544	1750000	3850	3	1	2	1	0	

	basement	hotwaterheating	airconditioning	parking
0	no	0	1	2
1	no	0	1	3
2	yes	0	0	2
3	yes	0	1	3
4	yes	0	1	2
...
540	yes	0	0	2
541	no	0	0	0
542	no	0	0	0
543	no	0	0	0
544	no	0	0	0


```
[545 rows x 11 columns]
0    furnished
1    furnished
2    semi-furnished
3    furnished
4    furnished
...
540   unfurnished
541   semi-furnished
542   unfurnished
543   furnished
```

```
In [5]: X = df[['area', 'bedrooms', 'bathrooms',
               'stories', 'parking']]
y = df['price']

In [6]: from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.4, random_state=101)
```

```
In [7]: from sklearn.linear_model import LinearRegression
lm = LinearRegression()
lm.fit(X_train, y_train)
```

```
Out[7]: LinearRegression()
```

```
In [8]: print(lm.intercept_)
-245089.43902304176
```

```
In [9]: coeff_df = pd.DataFrame(lm.coef_, X.columns, columns=['Coefficient'])
coeff_df
```

```
Out[9]:
```

	Coefficient
area	3.492829e+02
bedrooms	1.283724e+05
bathrooms	1.232385e+06
stories	5.085921e+05
parking	4.068205e+05

```
In [10]: predictions = lm.predict(X_test)
plt.scatter(y_test, predictions)
```


```
Out[10]: <matplotlib.collections.PathCollection at 0xid08ef3a790>
```





```
In [11]: sns.distplot((y_test-predictions),bins=50)
```

C:\Users\lenovo\anaconda3\lib\site-packages\seaborn\distributions.py:2619: FutureWarning: 'distplot' is a deprecated function. 'distplot' (a figure-level function with similar flexibility) or 'histplot' (an axes-level function for histograms).
warnings.warn(msg, FutureWarning)



```
In [ ]: from sklearn import metrics  
  
mse = mean_squared_error(y_test, predictions)  
print('Mean Squared Error:', mse)  
Mean Squared Error: 0.14268744338998152
```